

## DBMS\_REDEFINITION: Case Study for a Large Non-Partition Table to a Partition Table with Online Transactions occurring (Doc ID 1481558.1)

Modified: 02-Sep-2012    Type: BULLETIN

### In this Document

[Purpose](#)

[Details](#)

---

### APPLIES TO:

Oracle Server - Enterprise Edition - Version 11.2.0.2.0 and later  
Information in this document applies to any platform.

---

### PURPOSE

This Note has been written to provide some volume details and timings for the procedure to convert a normal table to a partition table using the DBMS\_REDEFINITION package.

The note will also include details of objects used by the process (Fast Refresh) and how these become populated during the procedure.

The procedure is broken down into the following sections:

1. Create the Partition Table structure required, known as the Interim table.
2. Execute DBMS\_REDEFINITION.can\_redef\_table...
3. Execute DBMS\_REDEFINITION.start\_redef\_table...
4. Execute DBMS\_REDEFINITION.sync\_interim\_table...
5. Execute DBMS\_REDEFINITION.finish\_redef\_table...
6. Drop the interim table.

---

### DETAILS

#### Worked example under 11.2.0.3

-- Initial setup of table to be partitioned.

```
CREATE TABLE unpar_table (  
  a NUMBER, y number,  
  name VARCHAR2(100), date_used date);  
  
alter table unpar_table ADD (CONSTRAINT unpar_table_pk PRIMARY KEY (a,y));  
  
-- load table with 1,000,000 rows  
  
begin  
  for i in 1 .. 1000  
  loop  
    for j in 1 .. 1000  
    loop  
      insert into unpar_table values ( i, j, dbms_random.random, sysdate-j );  
    end loop;  
  end loop;  
end;  
/  
commit;  
  
PL/SQL procedure successfully completed.
```

Elapsed: 00:01:56.90

Commit complete.

```
SQL> EXEC DBMS_STATS.gather_table_stats(user, 'unpar_table', cascade => TRUE);
```

```
SELECT  num_rows FROM user_tables WHERE table_name = 'UNPAR_TABLE';
      NUM_ROWS
-----
      1000000
```

Elapsed: 00:00:00.01

```
SQL> CREATE TABLE par_table (
      a NUMBER, y number,
      name VARCHAR2(100), date_used DATE)
PARTITION BY RANGE (date_used)
(PARTITION unpar_table_12 VALUES LESS THAN (TO_DATE('10/07/2012', 'DD/MM/YYYY')),
 PARTITION unpar_table_15 VALUES LESS THAN (TO_DATE('15/07/2012', 'DD/MM/YYYY')),
 PARTITION unpar_table_MX VALUES LESS THAN (MAXVALUE));
```

```
SQL> EXEC Dbms_Redefinition.can_redef_table(USER, 'unpar_table');
```

PL/SQL procedure successfully completed.

Elapsed: 00:00:00.07

```
-- This procedure (DBMS_REDEFINITION.start_redef_table) creates a materialized view based
on a CTAS, as we can see below with
-- the PREBUILT container table.
```

```
SQL> BEGIN
DBMS_REDEFINITION.start_redef_table(
uname => USER,
orig_table => 'unpar_table',
int_table => 'par_table');
END;
/
```

PL/SQL procedure successfully completed.

Elapsed: 00:00:06.69

```
select mview_name, container_name, build_mode from user_mviews;
MVIEW_NAME          CONTAINER_NAME          BUILD_MOD
-----
PAR_TABLE            PAR_TABLE                PREBUILT
```

Elapsed: 00:00:00.13

- Insert 1000 rows into the master table while the DBMS\_REDEF is active.
- This will use the mview log created by the DBMS\_REDEFINITION.start\_redef\_table.
- Check the MLOG\$\_<table name> table to confirm these online updates have been recorded.

```
SQL> begin
      for i in 1001 .. 1010
      loop
        for j in 1001 .. 1100
        loop
insert into unpar_table values ( i, j, dbms_random.random, sysdate-j );
        end loop;
      end loop;
    end;
  /
commit;
```

Elapsed: 00:00:00.24

```
SQL> select count(*) from MLOG$_UNPAR_TABLE;

COUNT(*)
-----
      1000
```

- Run the dbms\_redefinition.sync\_interim\_table to populate the new table structure (implements a MVIEW FAST REFRESH)
- with the online updates. This will purge the mview log of each record applied.
- This can be run many times and should be, before we do the Finish\_REDEF\_TABLE.

```
SQL> BEGIN
dbms_redefinition.sync_interim_table(
uname => USER,
orig_table => 'unpar_table',
int_table => 'par_table');
END;

PL/SQL procedure successfully completed.

Elapsed: 00:00:02.01

ALTER TABLE par_table ADD (CONSTRAINT par_table_pk2 PRIMARY KEY (a,y));

EXEC DBMS_STATS.gather_table_stats(USER, 'par_table', cascade => TRUE);
```

- Finish\_redef\_table swaps the table names so the interim table becomes the original table name.
- After completing this step, the original table is redefined with the attributes and data of the interim table.
- The original table is locked briefly during this procedure.

```
BEGIN
dbms_redefinition.finish_redef_table(
uname => USER,
orig_table => 'unpar_table',
int_table => 'par_table');
END;
/
```

- Note, both tables will now be synchronised.

```
select count(*) from par_table ;
COUNT(*)
-----
  1001000

select count(*) from unpar_table ;
COUNT(*)
-----
  1001000
```

- Dictionary views to confirm the change of structure of our original table "UNPAR\_TABLE".

```
SELECT partitioned FROM user_tables WHERE table_name = 'UNPAR_TABLE';

PAR
---
YES
```

```
SELECT partition_name, num_rows FROM user_tab_partitions WHERE table_name = 'UNPAR_TABLE';
```

PARTITION_NAME	NUM_ROWS
UNPAR_TABLE_12	980000
UNPAR_TABLE_15	5000
UNPAR_TABLE_MX	16000

-- At this point the Interim table can be dropped.

```
drop TABLE par_table cascade constraints;
```

#### Note for previous versions.

For tests done under 11.1.0.7 the timings for the sync\_interim\_table where the interim table is large (1,000,000 in our case) the timings for this are considerably longer. This was not reviewed in depth, but initial investigation looks like the mview refresh for DBMS\_REDEFINITION has been updated during the MERGE cycle.

example:

11.1.0.7  
SYNC of 1000 rows -- Elapsed: 00:01:36.30 (1.5 minutes)